

Introduction to Deep Learning

15. Image Augmentation, Fine Tuning, Style Transfer

STAT 157, Spring 2019, UC Berkeley

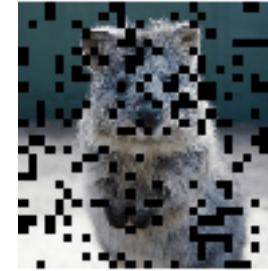
Alex Smola and Mu Li

courses.d2l.ai/berkeley-stat-157

Image Augmentation



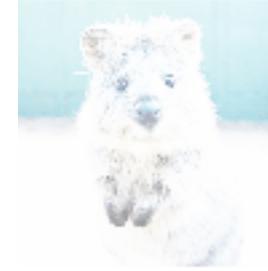
$p=1.0$



$\text{size_percent}=0.30$



$p=0.50$

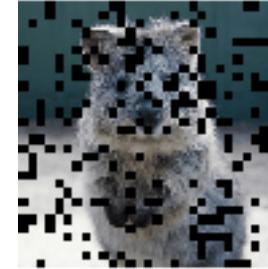


$\text{cutoff}=0.00$

Image Augmentation



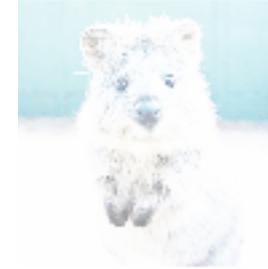
$p=1.0$



$\text{size_percent}=0.30$



$p=0.50$



$\text{cutoff}=0.00$

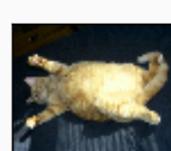
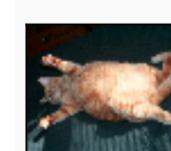
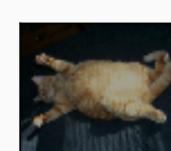
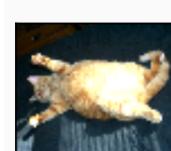
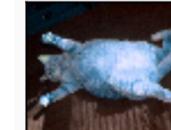
A Real Story at CES'19

- A startup found their demo, a smart vending machine that identifies what customers picked through a camera, didn't work because the showroom has
 - a different light temperature
 - light reflection from the desk
- They worked all night to re-collect data and train a new model, and ordered a tablecloth

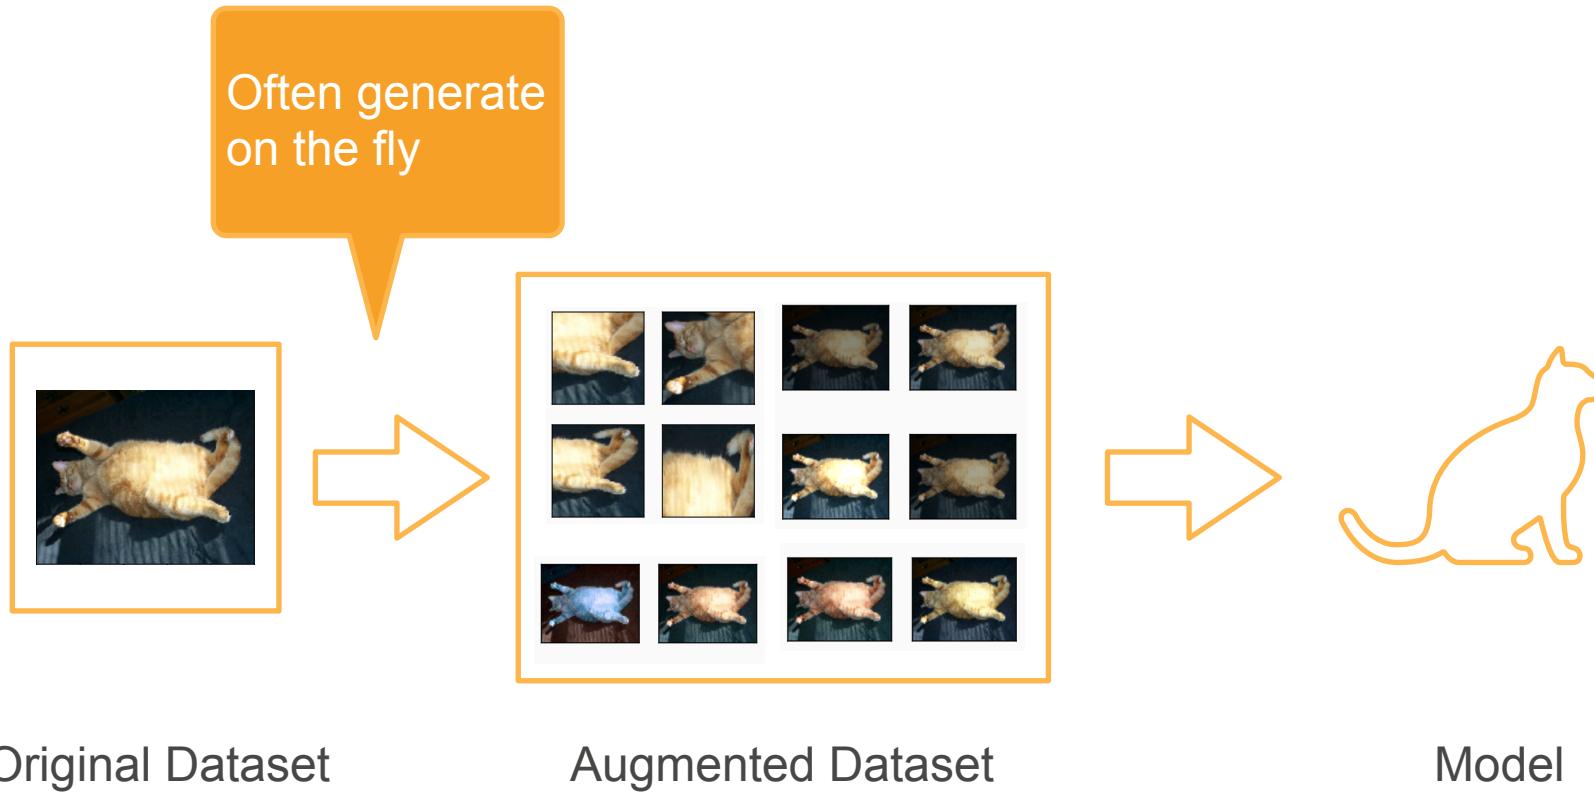


Data Augmentation

- Augment an existing dataset with more diversities
- Add various background noises into a speech
- Transform an image to several others by altering colors or/and changing shapes



Training with Augmented Data



Flip

- Left-right flip



- Top-bottom flip



- Not always makes sense



Crop

- Crop an area from the image and then resize it
 - A random width-height ratio (e.g. [3/4, 4/3])
 - A random area size (e.g. [8%, 100%])
 - A random position



Color

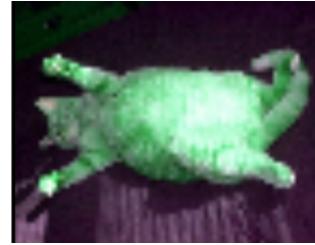
- Scale hue, saturation, and brightness (e.g. [0.5, 1.5])



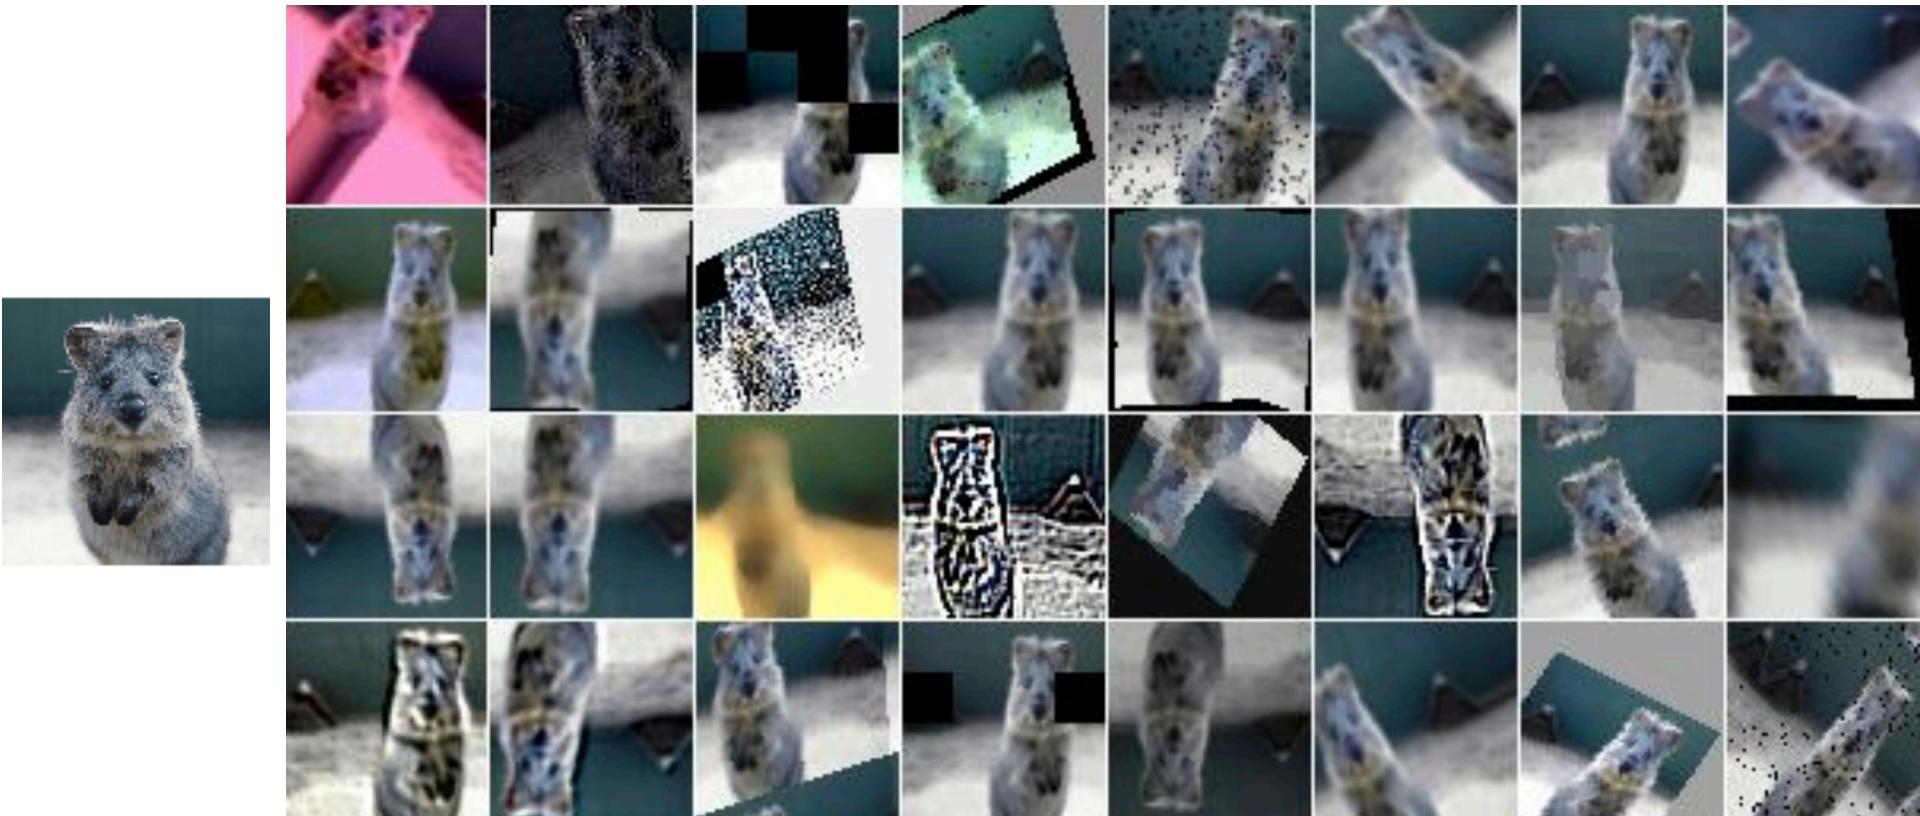
Brightness



Hue



Tens of Other Ways to Augment



Fine Tuning



Labelling a Dataset is Expensive

My dataset

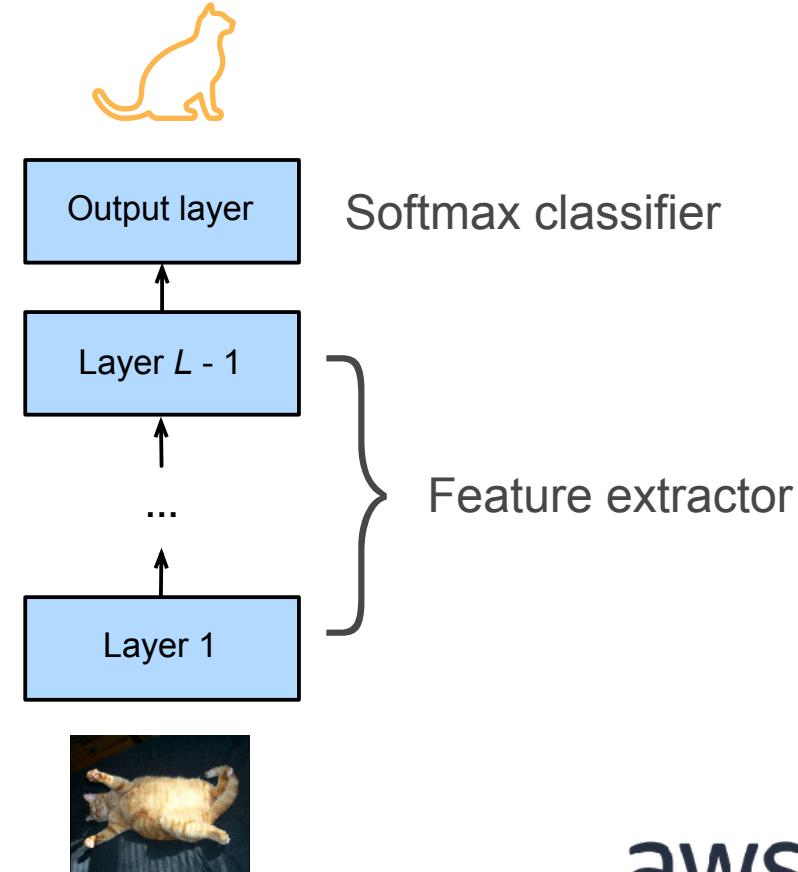


2 2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6 6

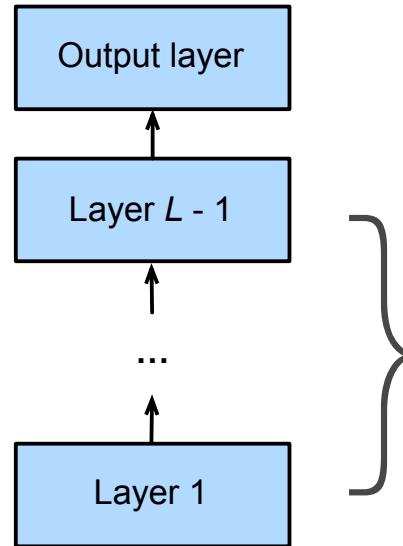
| | | | |
|------------|-------|-----|------|
| # examples | 1.2 M | 50K | 60 K |
| # classes | 1,000 | 100 | 10 |

Network Structure

- A neural network can be roughly partitioned into two parts
 - A feature extractor maps raw pixels into linearly separable features
 - A linear classifier to make decisions



Fine Tuning



Maybe not use the classifier parameters directly because labels are changed



Maybe also a good extractor for my dataset

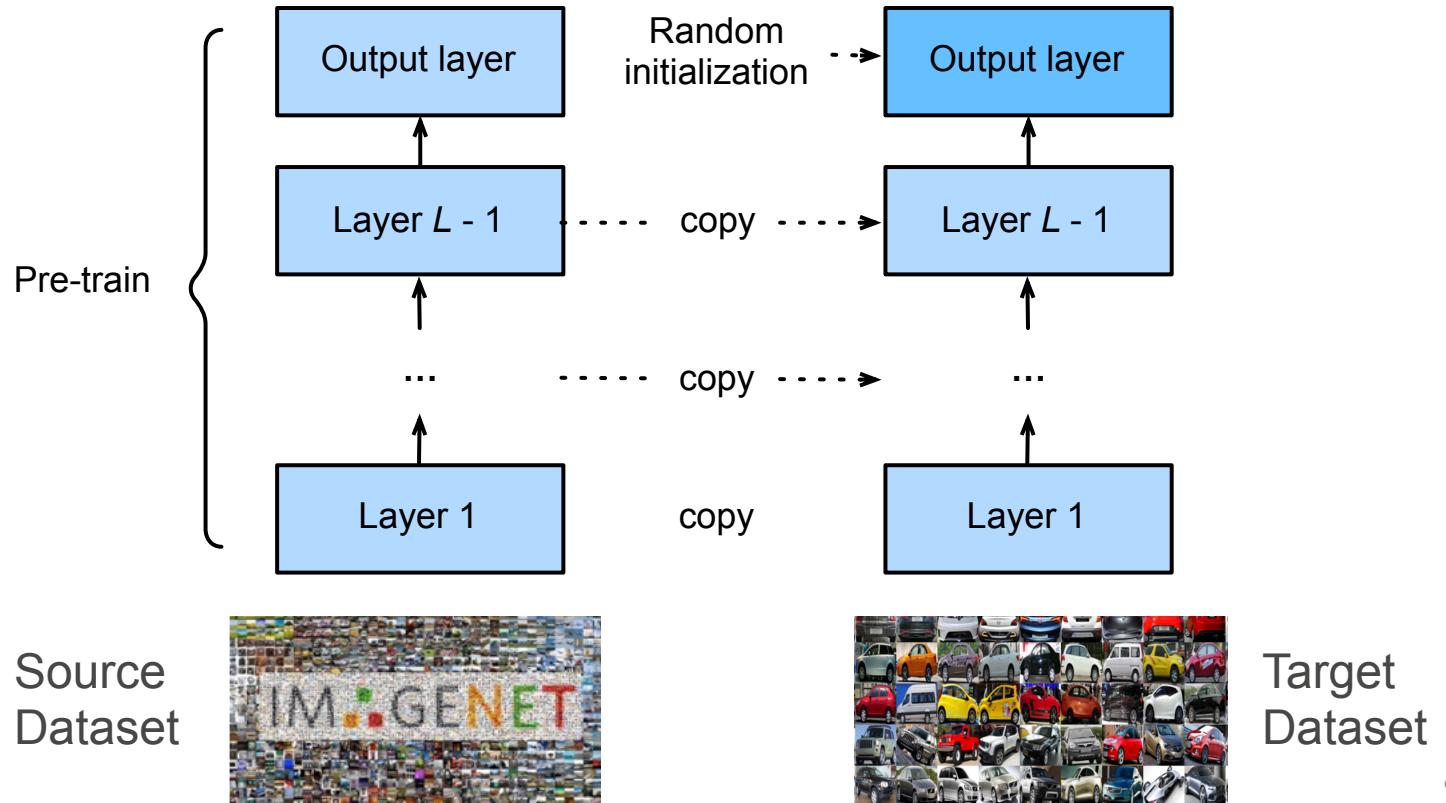
Source Dataset



Target Dataset



Weight Initialization for Fine Turning



Training

- Train on the target dataset as a normal training job, but with a strong regularization
 - Uses a small learning rate
 - Uses less epochs
- If source dataset is more complex than the target dataset, fine-tuning often leads to better quality models

Re-use Classifier Parameters

- The source dataset may contain some categories from the target datasets
- Use the according weight vectors from the pre-trained model during initialization



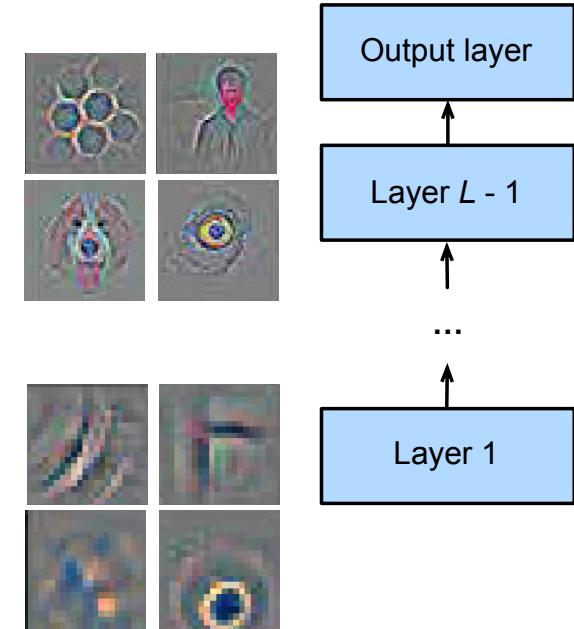
Racer, race car, racing car

A fast car that competes in races



Fix Some Layers

- Neural networks learn hierarchical feature representations
 - Low-level features are universal
 - High-level features are more related to objects in the dataset
- Fix the bottom layer parameters during fine tuning
 - Another strong regularizer



Style Transfer



+



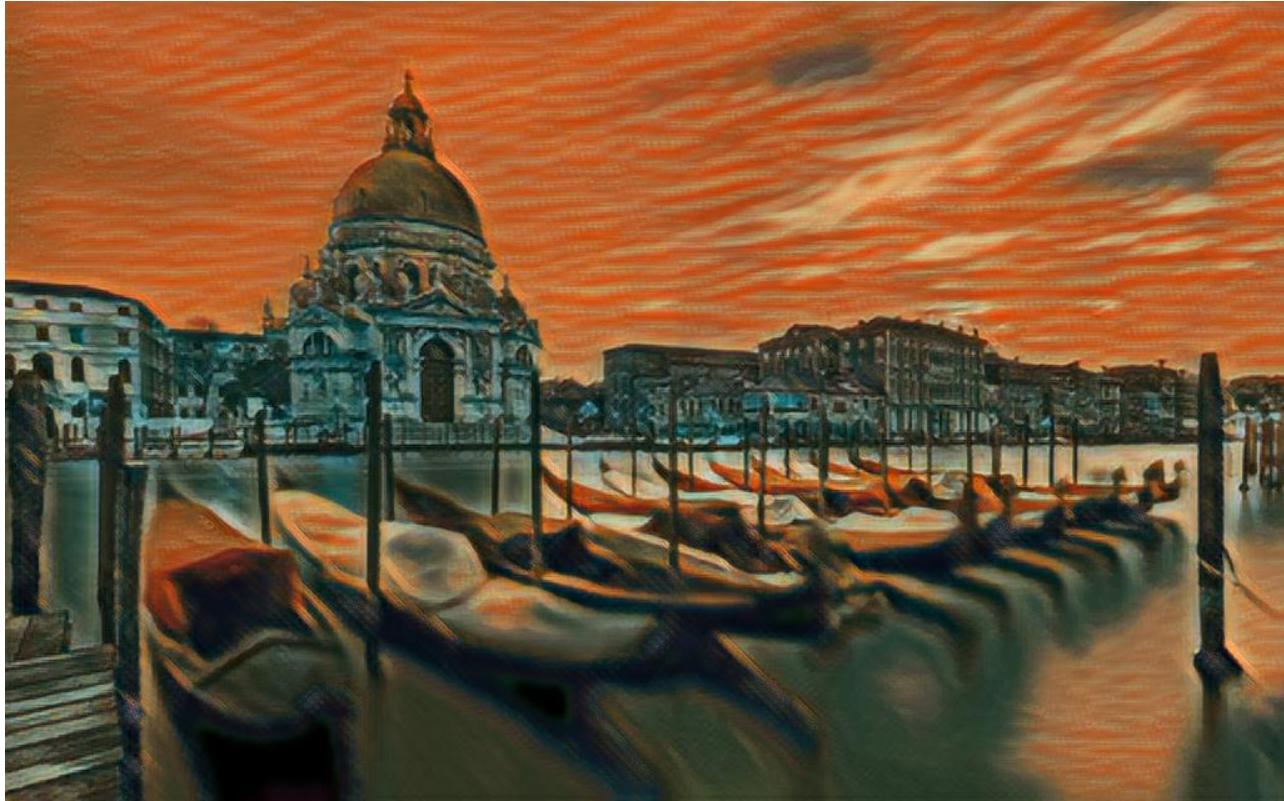
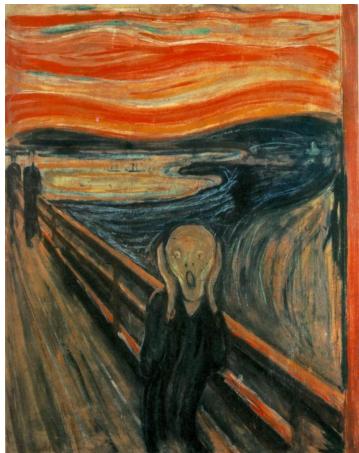


+





+





+





+





+





+



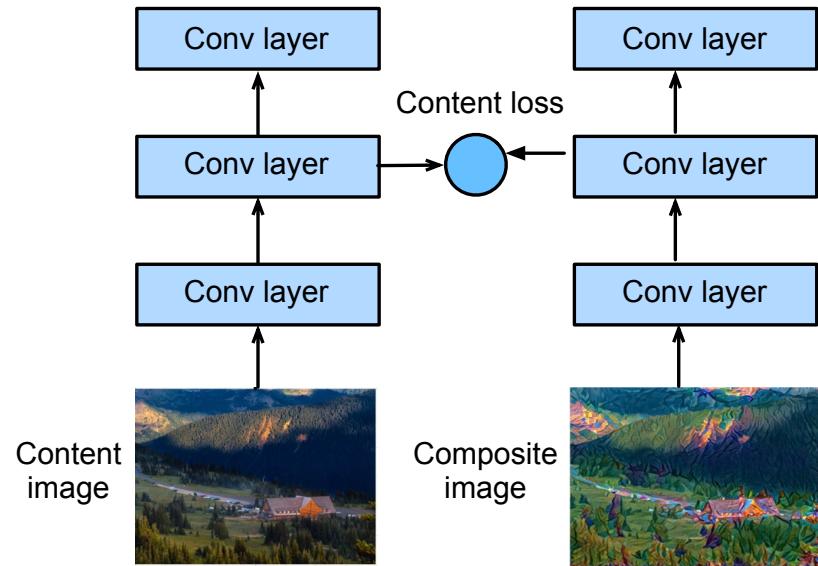
Neural Style

- Learn a composite image to match the contents from a content image and the styles from the style image

$$\arg \min_I w_1 \ell_{\text{content}}(I_{\text{content}}, I) + w_2 \ell_{\text{style}}(I_{\text{style}}, I) + w_3 \ell_{\text{noise}}(I)$$

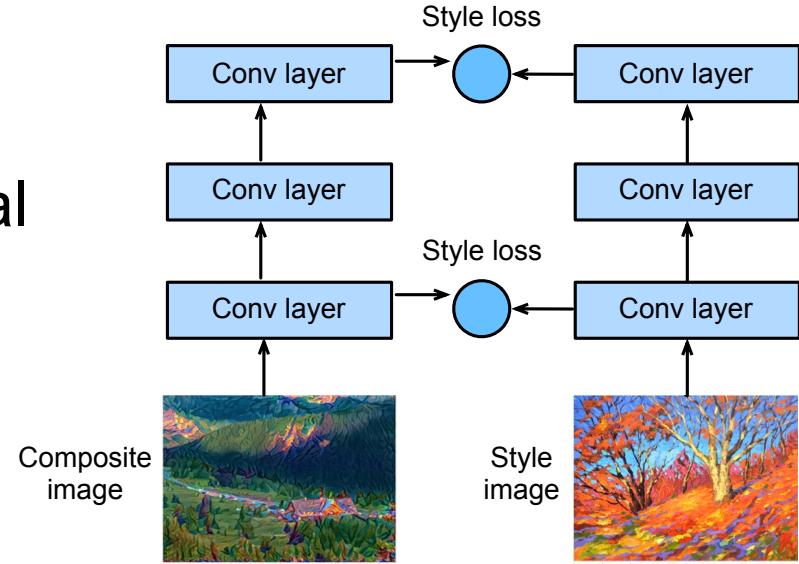
Content Loss

- Feed both content and composite images to the a CNN
- Compare internal layer outputs with a squared loss
 - Bottom layers matches details
 - Top layers matches global contents



Style Loss

- Gram matrix $G : G_{i,j}$ is the inner product between channel i and j
- Compare gram matrices of internal layer outputs by a squared loss
 - Bottom layers matches local styles
 - Top layers matches global styles



Noise Loss

- The learned composite image may have a lot of high-frequency noise
- Use total variation to de-noise

$$\sum_{i,j} |x_{i,j} - x_{i+1,j}| + |x_{i,j} - x_{i,j+1}|$$

Original image



Noisy image



Denoised image



Put All Things Together

$$\arg \min_I w_1 \ell_{\text{content}}(I_{\text{content}}, I) + w_2 \ell_{\text{style}}(I_{\text{style}}, I) + w_3 \ell_{\text{noise}}(I)$$

